

Deep V-Learning for Pool-Based Active Learning

March 31, 2022

1 Introduction

Even though pool-based Active Learning (P-AL) is the more popular setting for Active Learning, it creates significant challenges when reinforcement learning is applied to it. The most prominent one is that both Q-Learning and Policy Gradient Methods require a fixed action space. In the P-AL setting the action space has the same size as the pool of unlabeled images to choose from (or a subsample thereof). Even for simple datasets P-AL methods require a sample size of >100 to work effectively, creating a large action space for the reinforcement learning agent.

1.1 Contribution

- First V-Learning Approach for P-AL
- Variable sample sizes to fit every dataset
- (Dataset and model agnostic state space) *Probably not the first*

2 Related Work

3 Background

3.1 V-Learning

V-Learning ([1] p.119) is closely related to Q-Learning. It estimates the value of states rather than the value of all actions given a state (Q-Learning [1] p.131). This requires the environment to provide possible future states $s_{t+1} \in S_{t+1}$ given any state s_t . A policy π_v typically chooses the most promising future state $\operatorname{argmax} V(s_{t+1})$.

Temporal difference (TD) learning for a neural network parametrized by θ gives us a nearly identical update formula compared to Q-Learning

$$\theta \leftarrow \theta + \eta \left(r_t + \gamma \hat{V}_\theta(s_{t+1}) \right) \quad (1)$$

3.2 Active Learning as a V-Learning MDP

We change the usual MDP formulation from predecessor literature, which, based on a presented sample of unlabeled samples of size k , defines a state space $S := \mathbb{R}^{b \times k \times f}$ with features f and consequently an action space of $A := \{1, \dots, k\}^b$. Where b is the batch dimension, which was made explicit for clarity.

We instead define the state space to be 1-dimensional and use the batch dimension b for a variable sample size $S := \mathbb{R}^{b \times f}$. V-Learning does not define an action space, since the actions are implicit as chosen by the policy π_v .

This formulation poses a problem when storing transitions, however. The starting state s_t intuitively will be the chosen datapoint s_{t,π_t} , but since we don't want to store the full follow-up state $s_{t+1} \in \mathbb{R}^{b \times f}$ we use the average over the batch-dimension. This serves as a proxy of the impact of the implicit action $\pi_t := \pi(V(s_t))$ on the follow-up state, which is usually covered by the full state s_{t+1} . This results in a stored transition $\phi_t := \{s_{t,a_t}, r_t, \bar{s}_{t+1}, d_t\}$ where \bar{s}_{t+1} is the average of the follow-up state and d_t indicates whether a terminal state was reached.

4 Methodology

4.1 State Space

$$S := \mathbb{R}^{b \times (3+|Z|)} := [\text{BvsSB}, \text{Entropy}, \text{F1}, \text{Z}]^b \quad (2)$$

4.2 Model and Reinforcement Learning

We use a MLP with a single output neuron $f_\theta : \mathbb{R}^{3+|Z|} \rightarrow \mathbb{R}$ as agent network. As stated before, the sample of presented datapoints uses the batch dimension for sample size, resulting in b point-wise predictions of the agent.

For fitting the agent we use a target network f_{θ^-} [3], an n-step return ([1] p. 142) with $n = 3$ and prioritized experience replay [2].

5 EXPERIMENTAL: Dueling Networks for dynamic Action Spaces

$$\begin{aligned} \mathcal{S} &:= \mathbb{R}^{b \times k \times 3 + |Z|} ; \text{ with } b = 1 & \mathcal{S} &:= \{\mathcal{T}, \mathcal{C}\} := \{\mathbb{R}^{k \times 3}, \mathbb{R}^{|Z|}\} \\ \mathcal{A} &:= \{0, \dots, k\}^b & t &:= (s_t, r_t, \bar{s}_{t+1}, d_t) \\ t &:= (s_t, a_t, r_t, s_{t+1}, d_t) & \text{with : } s_t &:= \{t_{a_t}, c_t\} \\ & & & : \bar{s}_{t+1} := \{\bar{t}, c_t\} \\ Q^\pi(s, a) &:= \mathbb{E}[s_t, a_t, \pi] & Q^\pi(s, a) &:= \mathbb{E}[s_t, a_t, \pi] \\ V^\pi(s) &:= \mathbb{E}_{a \sim \pi(s)}[Q^\pi(s, a)] & V^\pi(s) &:= \mathbb{E}_{a \sim \pi(s)}[Q^\pi(s, a)] \\ A^\pi(s, a) &:= Q^\pi(s, a) - V^\pi(s) & A^\pi(s, a) &:= Q^\pi(s, a) - V^\pi(s) \end{aligned}$$

References

- [1] Andrew G. Barto Richard S. Sutton. *Reinforcement Learning: An Introduction*. MIT Press, Massachusetts, 2 edition, 2020.
- [2] Tom Schaul, John Quan, Ioannis Antonoglou, and David Silver. Prioritized experience replay. *arXiv preprint arXiv:1511.05952*, 2015.
- [3] Hado Van Hasselt, Arthur Guez, and David Silver. Deep reinforcement learning with double q-learning. In *Proceedings of the AAAI conference on artificial intelligence*, volume 30, 2016.