Research Presentation

Active Learning with V-Learning

Thorben Werner

March 31, 2022

Information Systems and Machine Learning Lab (ISMLL)
Institute for Computer Science
University of Hildesheim

Pool-Based Active Learning

$\mathcal{L} \in \mathbb{R}^{\lambda \times m}$ Labeled Set

$\mathcal{U} \in \mathbb{R}^{\mu \times m}$ Unlabeled Set

$\phi_\theta := \mathbb{R}^m \to \mathbb{R}^c$ Classifier

$\mathcal{K} \in \mathbb{R}^{k \times m}$ Unlabeled Sample

$\psi := \mathbb{R}^{k \times m} \to \mathbb{R}^k$ Active Learning Heuristic

$\pi_\psi := \mathit{argmax}\ \psi(\mathcal{S})$ Active Learning Policy

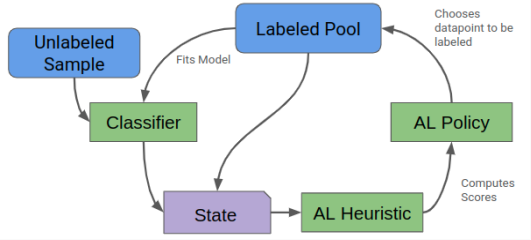

Active Learning Cycle

State Space:

1. Internal state of the environment: content of labeled pool $\mathcal{L}$, state of the classifier θ, remaining budget, current F1-Score, etc. $\rightarrow \mathbb{R}^a$
2. Information about the unlabeled sample $\mathcal{K}$: Output of the classifier $\phi_\theta(k_t)$, the datapoints themselves, other metrics, etc. $\rightarrow \mathbb{R}^{k \times b}$

Results in a flattened state space of $\mathcal{S} \in \mathbb{R}^{a+k \times b}$

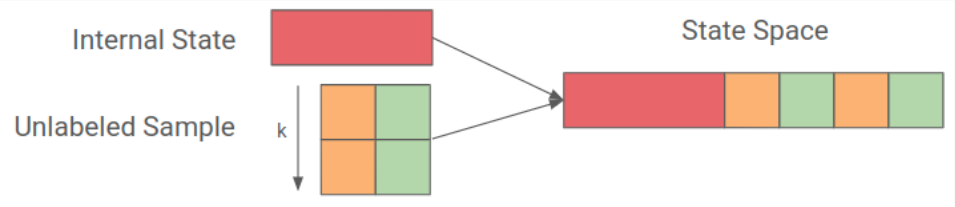

Active Learning with Q-Learning

$\mathcal{A} \in [0, \ldots, k]$ Action Space

$\mathcal{R} := \mathcal{S} \times \mathcal{A} \to \mathbb{R}$ Reward Function

$\tau := \{\mathcal{S}, \mathcal{A}, \mathcal{S}, \mathcal{R}, \mathbb{R}\}$ Transition

Problems:

- Fixed Sample size
- Expensive Transitions
- Same actions in different places

Sample Size $k = 2$

The same datapoint can appear in multiple places in the input

Both output nodes essentially learn the same function since the datapoints are sampled randomly and independently.

A permutation invariant network can fix the problem and create a ranking of actions

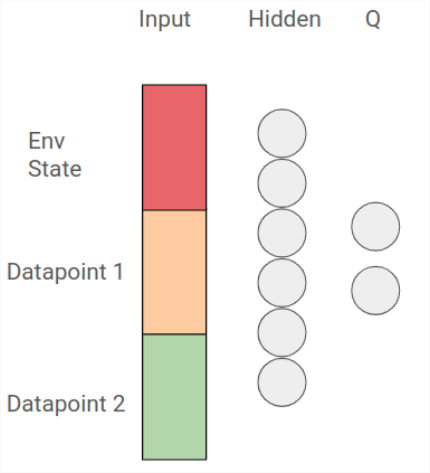

The Bellman Target

Q-Learning Target: What is the expected discounted improvement in F1-Score under the current policy when we choose a given datapoint?

The Bellman target does not rank actions, but tries to estimate their independent, global value under the current policy

Reinforcement Learning

$\mathcal{L}_Q := \hat{\pi}(s_t, a_t) - r_t + \gamma \max_a(\pi(s_{t+1}))$

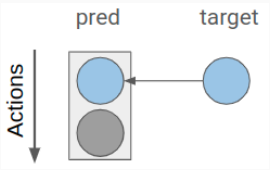

Ranking (Triplet Loss)

$\mathcal{L}_{triplet} := max(0, m + d(z_a, z_p) - d(z_a, z_n))$

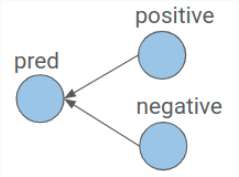

V-Learning vs Q-Learning

TODO

Adaptations

- We use the batch dimension for representing the sample size k
- No action space needed

$\mathcal{S} \in \mathbb{R}^{\sigma}$ State Space

$\mathcal{R} := \mathcal{S} \to \mathbb{R}$ Reward Function

$V_\theta := \mathcal{S} \to \mathbb{R}$ Agent Network

Example:

$$\text{State is generated} : s \in \mathbb{R}^{b \times \sigma}$$

$$\text{Agents makes prediction} : v = V_\theta(s)$$

$$\textbf{Policy selects a point} : a = \underset{b}{\text{argmax}}\, v_b$$

$$\text{Action is applied} : s' = \text{env}(a)$$

$$\text{Reward is observed} : r = \mathcal{R}(s')$$

$$\text{Transition is stored} : \tau = \{s_a, r, \bar{s}', \text{done}\}$$

$$\textbf{with} : \bar{s}' = \frac{1}{b} \sum_{i=0}^{b} s'_i$$

Storing Transitions in Memory

Q-Learning:

$$O(\tau) = s \times a \times s' \times r \times \mathbb{R}$$
$$O(\tau) = k^2 \times \sigma^2 + 3$$

V-Learning:

$$O(\tau) = s_a \times a \times \bar{s}' \times r \times \mathbb{R}$$
$$O(\tau) = 2 \times \sigma^2 + 3$$