

Research Presentation

Active Learning with V-Learning

Thorben Werner

March 31, 2022

Information Systems and Machine Learning Lab (ISMLL)
Institute for Computer Science
University of Hildesheim

Methodology

$\mathcal{L} \in \mathbb{R}^{\lambda \times m}$ Labeled Set

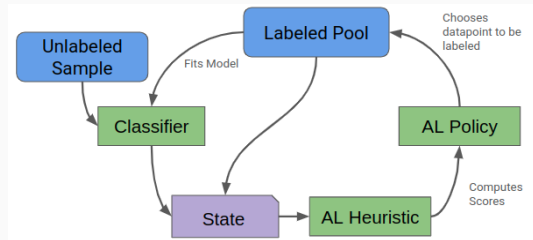
$\mathcal{U} \in \mathbb{R}^{\mu \times m}$ Unlabeled Set

$\phi_{\theta} := \mathbb{R}^m \rightarrow \mathbb{R}^c$ Classifier

$\mathcal{K} \in \mathbb{R}^{k \times m}$ Unlabeled Sample

$\psi := \mathbb{R}^{k \times m} \rightarrow \mathbb{R}^k$ Active Learning Heuristic

$\pi_{\psi} := \operatorname{argmax} \psi(\mathcal{S})$ Active Learning Policy




Active Learning Cycle

State Space:

1. Internal state of the environment: content of labeled pool \mathcal{L} , state of the classifier θ , remaining budget, current F1-Score, etc. $\rightarrow \mathbb{R}^i$
2. Information about the unlabeled sample \mathcal{K} : Output of the classifier $\phi_{\theta}(k_t)$, the datapoints themselves, other metrics, etc. $\rightarrow \mathbb{R}^{k \times j}$

Results in a flattened state space of $\mathcal{S} \in \mathbb{R}^{i+k \times j}$



$\mathcal{S} \in \mathbb{R}^{i+k \times j}$ State Space

$\mathcal{A} \in [0, \dots, k]$ Action Space

$\mathcal{R} := \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ Reward Function

$\tau := \{\mathcal{S}, \mathcal{A}, \mathcal{S}, \mathcal{R}, \mathbb{R}\}$ Transition

Problems:

- Fixed Sample size
- Expensive Transitions
- Same actions in different places

Example: Pick a Card

Consider a cardgame where the agent needs to collect cards with high value

Each card is worth a certain number of points, and the goal is to collect the maximum amount of points

At each iteration the agent is presented two cards to choose from

Q-Learning is not the correct choice of agent here




Sample Size $k = 2$

The same datapoint can appear in multiple places in the input

Both output nodes essentially learn the same function since the datapoints are sampled randomly and independently.

A permutation invariant network can fix the problem and create a ranking of actions

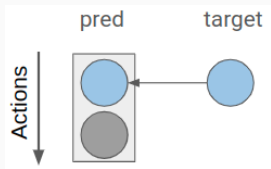


Q-Learning Target: What is the expected discounted improvement in F1-Score under the current policy when we choose a given datapoint?

The Bellman target does not rank actions, but tries to estimate their independent, global value under the current policy

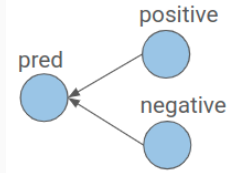
Reinforcement Learning

$$\mathcal{L}_Q := \hat{\pi}(s_t, a_t) - r_t + \gamma \max_a (\pi(s_{t+1}))$$



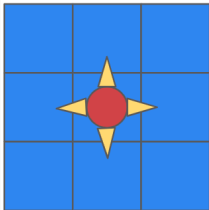
Ranking (Triplet Loss)

$$\mathcal{L}_{triplet} := \max(0, m + d(z_a, z_p) - d(z_a, z_n))$$



Q-Learning:

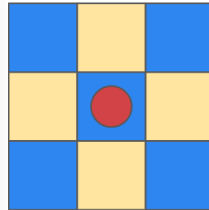
$$Q(s_t) = \mathbb{R}^4$$



$$Q(s_t) = \begin{bmatrix} V(s_{t+1}(\text{up})) \\ V(s_{t+1}(\text{down})) \\ V(s_{t+1}(\text{left})) \\ V(s_{t+1}(\text{right})) \end{bmatrix}$$

V-Learning:

$$V(s_{t+1}(a)) = \mathbb{R} \quad \text{with} \quad a = [1, \dots, 4]$$



$$Q(s_t, a) = V(s_{t+1}(a))$$

Adaptations

- We use the batch dimension ϕ for representing the sample size k
- No action space needed

$\mathcal{S} \in \mathbb{R}^{i+j}$ State Space

$\mathcal{R} := \mathcal{S} \rightarrow \mathbb{R}$ Reward Function

$V_{\theta} := \mathcal{S} \rightarrow \mathbb{R}$ Agent Network

Example:

State is generated : $s \in \mathbb{R}^{\phi \times (i+j)}$

Agents makes prediction : $v = V_{\theta}(s)$

Policy selects a point : $a = \underset{\phi}{\operatorname{argmax}} v_{\phi}$

Action is applied : $s' = \operatorname{env}(a)$

Reward is observed : $r = \mathcal{R}(s')$

Transition is stored : $\tau = \{s_a, r, \bar{s}', \text{done}\}$

$$\textbf{with} : \bar{s}' = \frac{1}{\phi} \sum_{p=0}^{\phi} s'_p$$

Q-Learning:

$$O(\tau) = s \times a \times s' \times r \times \mathbb{R}$$

$$O(\tau) = k^2 \times \sigma^2 + 3$$

V-Learning:

$$O(\tau) = s_a \times a \times \bar{s}' \times r \times \mathbb{R}$$

$$O(\tau) = 2 \times \sigma^2 + 3$$

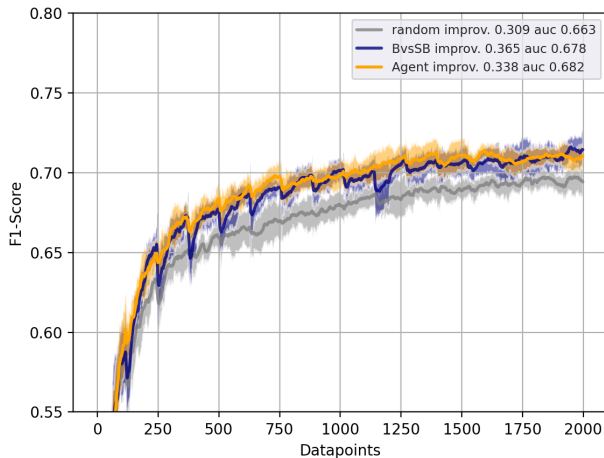
Current Status

V-Learning

CIFAR-10 (Custom
Embedding)

Samplesize: 20

Interactions: 2.4M




- Extremely hard to train a sequence of 2000 interactions
- Seemingly sensible state spaces collapse the performance
- Agent is overly sensitive

To test different agents and perform a quasi-gridsearch in feasible time I created a mock environment

Internal state is a fictional F1-Score

Each datapoint has a random "quality" value that governs its mock Entropy and BvsSB scores

A chosen sample add its "quality" value to the internal F1-Score



Internal state

$$\theta = 0 \quad // \text{current F1 score}$$

$$\mu = \text{Uniform}(0, 0.9)^k \quad // \text{current qualities}$$

State

$$\text{BvsSB} = \mu + \mathcal{N}(0, 0.2)$$

$$\text{entropy} = 2 + (\mu - 0.6) * 2 + \mathcal{N}(0, 1)$$

$$s = [\theta, \text{entropy}, \text{BvsSB}]$$

Step

$$\theta = \theta + \mu_a + \mathcal{N}(0, 0.5)$$

$$\mu = \text{Uniform}(0, 0.9)^k$$

Hand-picked combinations of "Budget", "Interactions", "LR", etc.

3-fold cross validation

Measuring regret with respect to BvsSB baseline

Preliminary results:

- N-Steps have seemingly no benefit
- Budget is more important than noise in the environment

